

Sensing-Fault-Aware Reinforcement Learning for Droplet Routing on Micro-Electrode-Dot-Array Biochips

Dr. Ankur Gupta

Abhay Raj Yadav

Hisham Siddiqui

Krishna Pansari

Abstract

Prior reinforcement-learning routers for MEDA biochips are fault-aware for the fluidic layer but assume a fault-free sensing layer, trusting capacitive health maps that can themselves suffer noise, spoofing, dropout, and delay. This interim report documents the learnability milestone of a sensing-fault-aware alternative: a Double-DQN router that receives a per-cell residual-based confidence signal alongside reported health and learns how much to trust each reading, evaluated in a fully digital simulator with architecturally decoupled fluidic and sensing fault channels. We report a failed first full-budget run honestly — sparse rewards and vanilla DQN dynamics produced baseline loss divergence (peaking at 12441) against aware-agent value collapse — and the minimal, identically applied fixes that restored learning: potential-based shaping on reported positions, Double-DQN target, and Huber loss. Gated clean-sensing pilots over 1200 episodes give GREEN (confidence-aware probe plateau 0.42→0.56; baseline churning near 0.40 with converged losses), supported by bounded TD losses, healthy Q-spread, and verified greedy evaluation. No fault-robustness gap is claimed: faulty-phase training and the four-experiment battery (parity, fault sweep, ablation, co-occurrence) are the wired-up next milestone.

1 Introduction

Digital microfluidic biochips (DMFBs) manipulate discrete droplets across an array of actuated electrodes to execute biochemical assays. Microelectrode-dot-array (MEDA) architectures extend this paradigm with fine-grained, freely addressable microcells, enabling eight-directional droplet transport rather than the Manhattan routing of conventional electrode arrays [1]. Routing droplets efficiently around faulty cells is therefore a sequential decision-making problem, and recent work has applied reinforcement learning (RL) to it with promising results.

The strongest such line, exemplified by prior deep RL routing [1], trains RL routers that are *fault-aware for the fluidic layer*: the agent observes a real-time health map of the array and learns to route around dead or degraded electrodes. This report starts from a gap in that line of work: the sensing layer itself is assumed fault-free. The health map the agent trusts is produced by capacitive sensing circuitry that can itself suffer noise, spoofed readings, signal dropout, and latency — and a fluidic fault masked by a simultaneous sensing fault is precisely the case the trust-the-map agent handles worst. To our knowledge, no prior RL routing work models sensing faults, let alone the co-occurrence of fluidic and sensing faults on the same cell.

Our contribution is a sensing-fault-aware router: alongside the (possibly corrupted) health reading, the agent receives a per-cell *confidence signal* derived from the residual between the commanded (expected) droplet position and the reported position, smoothed over a rolling window. The agent *learns* how much to trust each reading; there is no hard-coded cutoff. The baseline conditions differ only in the absence of this signal, so any performance gap is attributable to learned distrust rather than capacity. All experiments are fully digital: array, droplets, faults, and sensors are simulated, with fluidic and sensing fault channels architecturally decoupled.

This is an *interim* report. Its scope is the learnability milestone that precedes the headline comparison: we show (i) that the first full-budget training run failed and why, via logged evidence; (ii) the minimal, claim-preserving fixes; and (iii) gated pilot evidence that both agents learn under the fixed setup over 1200 clean-sensing episodes. The fault-robustness comparison (clean parity, sensing-fault degradation curves, confidence ablation, co-occurrence) awaits the full retrain and is explicitly out of scope here. Section 3 describes the methodology, Section 4 the implementation and the 1200-episode results, and Section 5 the path to the full evaluation.

2 Literature

2.1 MEDA architecture

Micro-electrode-dot-array (MEDA) biochips, first introduced in [2], place a CMOS control and sensing layer beneath every microelectrode, giving per-cell actuation and capacitive readout at fine pitch. Figure 1 shows the scale of this integration: a packaged chip (top left), the excitation/readout circuits tiled across the array (a), the sensing-pixel schematic with its transimpedance-amplifier readout (b), and the electrode profile and cross-section (c). This density enables free-direction routing and fine-grained droplet control — and a fault space strictly broader than that of conventional EWOD chips, surveyed below. Faults divide throughout into *catastrophic (hard)* faults, causing total malfunction, and *parametric (soft)* faults, degrading performance until a tolerance threshold is crossed.




Figure 1: MEDA sensing integration: packaged chip with excitation/readout circuits (a), sensing-pixel schematic with transimpedance-amplifier readout (b), and electrode profile/cross-section (c). Every micro-cell carries its own sense circuitry — the same circuitry whose failures motivate this work.

2.2 Error modes and faults in MEDA biochips

Subsections below are ordered by descending volume of published research on each layer, with sensing-circuit faults last: they are the least modeled as a first-class failure mode, and this work’s target. Original letter IDs are kept as stable tags.

2.2.1 Catastrophic microelectrode-level faults

Three hard failure types, each total for the affected cell:

- **Dielectric breakdown** — sustained high actuation voltage punches through the dielectric layer, shorting droplet to electrode; electrolysis sets in and transport halts. On MEDA the damage stays localized (one microelectrode of many) but still slows or stops the droplet, and only the fluidic domain is affected.
- **Shorts between adjacent electrodes** — etching defects fuse two neighbors into one electrical node; a droplet straddling them can no longer be actuated in parts, impeding motion and reducing droplet volume.

- **Open/broken interconnects** — a break in the metal between an electrode and its activation circuitry leaves the electrode permanently unfireable (stuck-at); droplets depending on it freeze or deform, varying shape, size, and volume.

Existing methods: a deep RL router is trained in [1] on real-time health information to bypass degraded electrodes, cutting clock cycles and execution time against baselines on COVID-19 testing protocols, with validation on fabricated prototype chips (advantage: online adaptation with hardware validation); multiple simultaneous array faults are detected in [3] by routing test droplets in clockwise/anticlockwise pebble-traversal patterns (Borderline plus Inner traversal with traversal-time computation), then repaired by Advanced Module Sequence Graph reconfiguration that swaps defective cells with spares so the chip performs like new (advantage: strong multi-fault coverage plus repair); modules are placed in [4] by a DQN-based placement method (RLPM) that minimizes chip area while maximizing a reliability score, paired with A*-seeded collision-free routing (CAMR) that stalls lower-priority routes out of collision zones, cutting area use $\sim 28.6\%$ without compromising reliability (advantage: prevention plus assay-visible routing). *Drawbacks:* the deep RL router in [1] needs offline training plus an accurate simulator, is single-agent, and assumes its health sensing is fault-free; the pebble-traversal scheme in [3] is offline only (whole chip idle, traversal time scales with size, no mid-assay repair); synthesis-time placement in [4] cannot react to field faults.

2.2.2 Parametric and degradation faults

Three degradation mechanisms, all cumulative:

- **Charge trapping (electrode wear)** — every charge/discharge cycle piles irreversible charge near the electrodes, weakening actuation force until droplets stick, split astray, or stop; complete degradation is documented after roughly 200 charge cycles, and excessive actuation pushes the same mechanism over the edge into dielectric breakdown, making wear and breakdown one continuum.
- **Fabrication-tolerance deviations** — insulator thickness and electrode geometry drifting outside spec, shifting actuation behavior cell to cell from birth.
- **Droplet volume loss** — underlying defects bleed droplet volume, measurable through oscillation and impedance shifts, silently skewing every downstream volumetric operation.

Existing methods: explicit robustness to charge-trapping degradation is shown in [1] (advantage: handles the most common wear mechanism, validated on prototypes); degradation-aware routing is extended in [5] to concurrent multi-agent transport: one cooperative agent per droplet, trained in an open-sourced PettingZoo simulator, needing no external CCD camera, and demonstrated on a fabricated chip running serial dilution — with double DQN, PPO, and ACER compared across centralized, decentralized, and parameter-sharing schemes (advantage: parallelism without external sensing hardware). *Drawbacks:* same simulator/single-agent and sensing-assumption limits as above; the multi-agent extension in [5] pays coordination overhead and training complexity that grows with droplet count.

2.2.3 Daisy-chain and control-path faults

Three failure points along the chain:

- **Stuck control registers** — D flip-flops under the MCs freeze (stuck-at blocks control-signal changes), so activation patterns cannot be written past them.
- **Opens between neighboring MCs** — a broken link stops data transmission down the chain entirely; unlike a dead electrode, nothing downstream is reachable.
- **Whole-chain break** — either of the above, unrecovered, discards the chip: with all MCs serially chained (three control pads driving 900 MCs in the FPLOC prototype), one failed cell severs control for everything behind it.

Existing methods: online diagnosis with functional actuation patterns in [6] repairs faulty cells mid-assay using data already shifting through the chain — no influence on the executing bioassay, no added degradation, no extra time (advantage: zero-downtime repair); enhanced built-in self-diagnosis and self-repair in [7] widens resolution and coverage over the earlier 2018 BISD scheme, adding a bypass multiplexer per MC plus JTAG-configured recovery that walks faulty chains cell by cell, distinguishing recoverable from unrecoverable MCs separately for activation versus sensing operations (advantage: finer, broader chain repair with per-operation recovery accounting). *Drawbacks:* both are scoped to the chain (blind to fluidic and sensing faults) and bounded by per-cell redundancy, with hardware overhead scaling in chip size.

2.2.4 System and bioassay-level consequences

No fault stays local; the cascade runs:

- **Stuck/misrouted droplets** — the direct surface symptom of any cell-level failure above.
- **Routing collisions** → **cross-contamination** — droplets meeting where they shouldn't, mixing samples.
- **Skewed dilution/mixing ratios** — stranded or split droplets corrupt volumetric steps.
- **Assay-timing violations** — every recovery attempt consumes protocol time until the schedule cascades; in point-of-care diagnostics the chain ends at erroneous readouts and misdiagnosis risk.

Existing methods: tests for parallel fluidic modules are scheduled in [8] (advantage: module-level coverage at lower test cost); collision-avoidance routing [4] and pebble-based reconfiguration [3] operate at this level (advantage: assay-visible reliability); contamination-aware cooperative MARL routing is pursued in [9], cutting cross-contamination events 47–58% against the state of the art while holding success rates and completion steps (advantage: measured contamination reduction). *Drawbacks:* symmetric-module testing in [8] depends on module symmetry and irregular layouts defeat it; synthesis-time methods cannot respond mid-assay; offline reconfiguration needs an idle chip; MARL coordination adds overhead.

2.2.5 MEDA-specific advanced-operation faults

Three classes with no conventional counterpart, in a literature that overwhelmingly assumes four-neighborhood movement and therefore never exercises diagonal paths:

- **Contact-line pinning during diagonal movement** — corner traversal the Manhattan model never exercises; undetectable by conventional test methods by definition.
- **Size-dependent velocity faults** — droplet speed varies with shape, volume, and direction, so timing assumptions from fixed-electrode models misfire.
- **Split/merge geometry faults** — morphing and splitting restricted to physically plausible patterns; layouts violating them fail in ways fixed-module testing never probes.

Existing methods: none surveyed target these directly. The verified speed lineage (67.5% solving-time reduction [10], 72.3% plus extended solvability in [11], via target-guidance constraints and lower-bound pruning) only indirectly helps, by making re-routing around faults real-time tractable. *Drawbacks:* speed is not diagnosis — without a fault model for these classes, faster solving just repeats blind decisions.

2.2.6 Security-induced (cyberphysical) faults

Five attack surfaces, all riding real interfaces:

- **Actuation tampering** — altered daisy-chain sequences corrupting outcomes while reporting success.
- **Sensor spoofing** — fabricated readings misleading recovery (the malicious twin of our spoof class).
- **Denial of service** — stalled time-critical protocols.

- **Sample/assay tampering** — corrupted inputs or protocols mid-run.
- **IP piracy** — design extraction through test/debug interfaces.

Existing methods: surveyed as concerns (e.g. the 2025 MEDA review[12] and recent sustainability/security surveys[13] flag microelectrode degradation, excessive data-shift operations, and MEDA security exposure); detection, prevention, and rectification techniques specifically for cyberphysical attacks on MEDA are developed in [14] — the one dedicated countermeasure line in scope-adjacent literature. *Drawbacks for our purposes:* attacker-induced faults share the sensing surface (notably spoofing) but require a threat model this work does not claim.

2.2.7 Sensing-circuit faults

Each microcell’s actuate/sense circuitry is $0.35\ \mu\text{m}$ CMOS replicated ~ 1800 times per chip, so it inherits CMOS defect densities directly: the fault-recovery literature documents intra-MC defects reaching the sensing circuits and builds backup readout multiplexers specifically so a faulty MC’s sensing result can be read through a neighbor. A stuck-at sense amplifier reporting “healthy” is phenomenologically our spoof class, and scan-chain read opens are our dropout class. Field exposure is worse: droplet sensing contributes up to 94% of all MC actuations [15], so the sense path wears first, while calibration drift, biofouling, and charge accumulation corrupt capacitive readings parametrically; the testing literature confirms sensing-circuit defects cause parametric *and* catastrophic failures that structural methods do not cover. Figure 2 shows the consequence: ground truth (left) versus the reported map (right) with spoofed cells (false positives), dropout cells (false negatives), and a misreported droplet position. *Existing methods:* component-level testing only (oscillation-based sensing tests, BIST built atop sensing circuits, sensitivity calibration) — no surveyed work models sensing faults as a first-class failure mode for routing; every sensing-map consumer treats the reported panel as ground truth. *Drawbacks of the status quo:* a fluidic fault masked by a simultaneous sensing fault on the same cell becomes an invisible trap, exactly where trust-the-map routing fails hardest. This layer is our work’s target.




Figure 2: True array state (ground truth, left) versus the reported sensing map under sensing faults (right): spoofed healthy cells, dropout cells, and a misreported droplet position. Consumers of the sensing map see only the right panel.

The pattern across layers is consistent: each line of work owns one layer, and no surveyed method co-diagnoses control-path, sensing-circuit, and fluidic faults together.

3 Methodology

Figure 3 shows the closed loop in full implementation detail: a MEDA simulator (array dynamics, droplet kinematics, fault injection, sensing model, rewards) exchanges observations and rewards with a DQN agent whose inputs are the array state, the sensing map, the confidence map, and the task goal. All sweeps are configuration-only (`fault_profiles.yaml`); no experiment requires code changes.




Figure 3: Agent-environment closed loop at implementation level: per-module simulator settings (array 10×10 , Discrete(8), decoupled faults, four sensing modes, shaped rewards) and the shared Double-DQN agent (matched encoder, Double target with Huber loss, continuous exploration); highlighted boxes are fix-cycle additions, identical across conditions except the confidence input.

3.1 Simulation environment

The array is a grid of microcells (default 10×10 , configurable), each holding ground-truth health $h \in [0, 1]$, a fault tag, cumulative wear, and droplet occupancy. The single droplet occupies one cell with a commanded target cell. Actuation offers eight directional moves (Discrete(8)): unlike the classic DMFB literature, which overwhelmingly assumes four-neighborhood (Manhattan) movement, MEDA’s free-direction routing makes diagonals first-class actions the agent selects among, and the simulator models them as such. Figure 4 shows the physical mechanism behind these moves. A move into a cell whose true health is at or below 0.3 silently fails (droplet stays put), modeling insufficient electrowetting force rather than a crash; out-of-bounds moves fail likewise. Diagonal moves validate only the destination cell, not the two orthogonal intermediaries (corner-cutting permitted): a deliberate modeling choice, flagged here because conventional DMFB routers typically forbid it, whose impact on learned policies is left for future ablations. Every actuated cell accrues wear (0.01/step), degrading health monotonically and progressing into dielectric breakdown at wear 1.0 — degradation and breakdown are one process, matching taxonomy (B). Array dimensions, voltage ranges (~ 20 – 40 V), and droplet volumes are chosen within MEDA-literature ranges and recorded in `array_default.yaml`; the dynamics abstract actuation to the health-threshold model above.




Figure 4: Electrowetting-on-dielectric actuation over a 2-D electrode array (cross-section inset): the physical mechanism behind the eight-directional moves the simulator models, including diagonals that classic four-neighborhood DMFB models omit.

3.2 Decoupled fault injection

Fluidic faults act *only* on ground-truth health: dielectric breakdown and stuck-at (catastrophic, permanent, $h=0$) and charge trapping (parametric, cumulative). Sensing faults act *only* on reported readings, never on ground truth: Gaussian noise ($\sigma=0.05\text{--}0.10$) on health; spoof reporting healthy (1.0) regardless of truth; dropout reading 0.0 (“no signal”) with zero-order hold on position; delay replaying lagged readings (position holds last reported — a documented freeze approximation). The four modes are catalogued in Figure 5. A per-episode scheduler samples the fluidic dead-cell set and the corrupted-sensor set *independently*, with the overlap fraction an explicit parameter (exact, not a lower bound): the regime the trust-the-map baseline handles worst. Fault density, type mix, and overlap live in `fault_profiles.yaml` (clean / low / medium / high / co-occurrence, plus a 0–50% density sweep); start and target cells are excluded from both sets.

3.3 Sensing-reliability (confidence) signal

The novel component. The droplet’s expected next position is derivable from the last commanded actuation; persistent deviation of the *reported* position implicates that cell’s sensor, independently of electrode health. Per cell we keep a rolling window (5) of residuals $r=\|\text{expected} - \text{reported}\|_2$ and report

$$c = \frac{1}{1 + \bar{r}} \in (0, 1], \quad (1)$$

so one-off noise washes out while persistent spoofing or freezing drives confidence down; all other cells relax toward 1 each step. The tracker update is skipped when the decision-time *reported* map already marks the destination dead (the residual then carries no position-sensor information) — a reported-only gate, never ground truth. The agent receives c as an input channel and *learns* the trust weighting; there is no hard-coded cutoff, keeping the baseline comparison fair. Figure 6 diagrams the full pipeline from capacitance computation to the confidence map.




Figure 5: The four simulated sensing-fault modes: Gaussian health noise, spoofed (false-healthy) readings, dropout with missed detections, and temporal delay. Each corrupts only the reported plane, never ground truth.

3.4 Agents

Baseline and confidence-aware agents are the *same* Double-DQN implementation (online argmax, target evaluation [16], Huber loss, Adam, target sync every 500 gradient steps) with identical hyperparameters ($lr=10^{-3}$, $\gamma=0.99$, buffer 20000, batch 64, $\epsilon: 1.0 \rightarrow 0.05$ over 800 episodes) and matched depth/width: a two-layer CNN




Figure 6: Simulated capacitive-sensing pipeline (left): capacitance computation with noise/spoof/dropout/delay injection yields the reported map, while the expected-vs-reported residual with temporal consistency yields the confidence map in $[0, 1]$ (right).

encoder with adaptive pooling plus MLP head. DQN is chosen deliberately: the prior deep RL routing lineage [1] is DRL-based, so the baseline must come from the same family for the comparison to mean anything; eight discrete moves suit value-based methods without continuous-control machinery; and single-agent scope defers multi-agent coordination overhead until the core sensing result is proven. The *only* model difference is input width (3 channels baseline: reported health, droplet, target; 4 aware: plus confidence), and the only environment difference is the confidence channel with its small suspect penalty. Rewards: +10 goal, −0.05/step, −1.0 for attempting a ground-truth-dead cell, −0.1 (aware only) for entering a sub-0.5-confidence cell. Dense guidance comes from potential-based shaping [17], $F = \gamma\Phi(s') - \Phi(s)$ with $\Phi = -k \cdot d_{\text{Chebyshev}}$, $k=0.2$ — policy-invariant, hence comparison-preserving — computed on *reported* positions, i.e. a function of the observation, consistent with the paper’s thesis.

3.5 Training curriculum and evaluation protocol

Phase A trains 1500 episodes under clean sensing (parity sanity check); phase B trains 1500 under a fixed medium sensing-fault rate so failures are encountered during learning. Phase 2 is *training*, not fine-tuning: network weights transfer across the boundary while the replay buffer is cleared, and one continuous epsilon schedule spans both phases. Both agents share paired curricula (identical reset-seed sequences) and paired greedy probes (every 250 episodes, $n=100$ for $\sim\pm 5\text{pp}$ selection noise); best-by-success checkpoints are saved and evaluation defaults to them, with fail-closed loading. Metrics are routing success rate, degradation versus sensing-fault density, false-trust rate (moves toward ground-truth-dead cells reported healthy), steps-to-goal, and greedy decision latency (warmed-up forward passes, reported with hardware as a ratio, not an absolute). The four-experiment battery (clean parity; 0–50% fault sweep; confidence ablation; co-occurrence) is defined in configuration; its numbers belong to the full report, not this interim one. Sweep densities are stress-test ranges, not claimed field rates.

4 Implementation and Results

4.1 Implementation map

The methodology of Section 3 is implemented in the public repository linked below, whose README maps every module; the load-bearing facts for this paper are three: baseline and aware agents are the same shared Double-DQN class differing only in input width, all sweeps are configuration-only, and thirty-six unit tests pin the reported-position, overlap, seeding, lockstep, gating, and shaping invariants. Figure 3 maps the same loop at implementation level: every fix-cycle addition is highlighted, and the two conditions share all unhighlighted boxes. The complete codebase, configurations, tests, and this report are publicly available at <https://github.com/arx13/MEDA-Sensing-Aware-RL>; run-generated artifacts (checkpoints, TensorBoard logs) are excluded from the repository and reproduced via the documented commands.

4.2 Setback: the first full-budget run failed to learn

The first 1500+1500-episode run, under the original sparse reward (+10 goal, −0.05/step) and vanilla DQN (MSE loss, max-bootstrap target), learned nothing: greedy probes peaked at 0.19 and decayed, with clean-sensing eval at 4.5% baseline / 11.5% aware — parity of failure. The loss curves split the diagnosis. The baseline diverged catastrophically (0.3→52 clean, spiking to 12441 under faults): rare +10 transitions arriving as massive TD errors inside a max-bootstrap feedback loop. The aware agent stayed numerically bounded (0.01–0.5) but its Q-values collapsed to the timeout value — a separate pathology, same outcome. The environment was exonerated first: a scripted Chebyshev-greedy policy scores 0.92 clean and 0.72 under medium sensing faults, so a learnable policy exists and the fault gap is real. The failure was learning configuration, and no fault-sweep numbers were drawn from these checkpoints: flat $\sim 5\%$ lines would have read as “no gap” when the truth was “no learning.”

4.3 Fixes

Three minimal, claim-preserving changes, applied identically to both agents:

- potential-based shaping ($k=0.2$, γ matched at 0.99) on *reported* positions, adding dense per-step gradient while preserving the optimal policy;
- Double-DQN target plus Huber loss, bounding the per-element gradient ($\delta=1$) and removing max-over-target feedback, with the learning rate untouched at 10^{-3} to isolate the effect;
- no architecture change, so any gain cannot be attributed to capacity.

Shaping signs and outlier-batch finiteness are unit-tested. Config-level diffs before/after (epsilon continuity, best-checkpoint selection, buffer clearing, $n=100$ probes) were verified identical across agents except for the observation channel — the earlier config audit found no confound.

4.4 Gated pilot evidence

A 300-episode clean pilot graded RED (final greedy probes 0.12/0.08), but against an $\varepsilon \approx 0.64$ schedule the probes measured a barely exploited policy, so the verdict was confounded rather than conclusive. Two sanity checks then passed: greedy evaluation is pure argmax with epsilon bypassed, and Q-spread over 20 states is healthy (median per-state Q-std ~ 0.25 , $\sim 7\times$ the random-init reference, argmax spread over 6–7 of 8 actions) — the encoder differentiates actions and TD targets flow, ruling out flat-Q collapse. The 900-episode extension (through full ε -decay) showed rising trajectories; the 1200-episode continuation, graded on trajectory, gave: aware 0.54 $\rightarrow$ 0.60 $\rightarrow$ 0.56 (GREEN: endpoint $\geq 40\%$, rise-then-plateau convergence) and baseline 0.46 $\rightarrow$ 0.42 $\rightarrow$ 0.36 (amber churn around ~ 0.40 with converged loss). Figure 7 records the full 12-probe trajectories against the 15/25/40% gate bands and the 0.92 scripted ceiling; Figure 8 the monotonically converging Huber losses (max 0.15 versus 12441 pre-fix); Figure 9 the behavior/exploration split explaining the early red; Figure 10 the Q-spread health check.




Figure 7: Greedy probe trajectories over 1200 clean episodes ($n=50$), both agents, with gate bands and the scripted-greedy ceiling.




Figure 8: TD-loss quarter means, both agents. Bounded, monotonically converging losses are the stability proof that separates this run from the failed one (pre-fix baseline peaked at 12441); without them, rising probes could be noise rather than learning.




Figure 9: Train-time (exploratory) success versus the epsilon schedule (baseline). Behavior succeeds on exploration luck while greedy probes lag, which is why premature probes mislead — this split is what made the 300-episode RED verdict confounded rather than conclusive.




Figure 10: Median per-state Q-value std over 20 states versus random-init reference. If Q were flat (all actions valued alike), the encoder would be the suspect and the fallback would change; the $\sim 7\times$ spread rules that out and pins the earlier failure on TD targets instead.

4.5 Limits of the current evidence

Three limits bound what is claimed. First, all runs above are clean-sensing phase-A equivalents: faulty-phase training and experiments 1–4 belong to the full retrain, and no fault-robustness gap is claimed here. Second, the baseline’s late probe churn ($\sim \pm 8\text{pp}$ around 0.40 at $n=50$) is unresolved detail — converged loss with a fluctuating greedy policy — which the full run’s $n=100$ probes and best-checkpoint selection are designed to settle. Third, pilot probes are single-seed (paired across agents, but one seed); multi-seed confidence intervals are deferred to the full evaluation. Fourth, all training and evaluation use the 10×10 default grid: whether the confidence signal survives at commercial array sizes, where prior deep RL routers operate at commercial scale [1], is untested and reserved for follow-up scaling studies. Full training (1500 clean plus 1500 faulty episodes, $n=100$ probes, best-checkpoint selection) is expected to raise clean-parity levels toward the scripted ceiling and to produce the faulty-phase gap that is this project’s headline; no numeric values are asserted before that run. A final exclusion, stated for the record: a composite project figure containing anticipated (not measured) fault-sweep bars was deliberately left out of this report — every number above traces to logged data.

5 Conclusion and Future Works

Prior RL routing for MEDA biochips is fault-aware for the fluidic layer and blind to the sensing layer: it treats a possibly corrupted sensing map as ground truth. This project relaxes that assumption with a router that receives a per-cell residual-based confidence signal and learns how much to trust each reading, evaluated in a fully digital simulator whose fluidic and sensing fault channels are architecturally decoupled. The matched baseline/aware design isolates any performance gap to learned distrust rather than capacity.

This interim report established the precondition for that comparison. The first full-budget run failed to learn under sparse rewards and vanilla DQN dynamics (baseline loss peaking at 12441); minimal, identically applied fixes — potential-based shaping on reported positions, Double-DQN target, Huber loss — restored stable learning with monotonically converging losses; and gated pilots gave GREEN on clean sensing (confidence-aware plateau $0.42 \rightarrow 0.56$; baseline churning around 0.40 with converged loss), clearing the way for the full retrain. No fault-robustness gap is claimed here: faulty-phase training and experiments 1–4 are the next milestone, with checkpoints, TensorBoard logs, and fail-closed evaluation already wired for it. The immediate next step is completing the full 1500+1500-episode retrain for both agents and running experiments 1–4 against it; beyond that, we plan systematic optimization — reward-scale and shaping-coefficient studies, hyperparameter sweeps around the stabilized Double-DQN, and the MLP ablation if encoder capacity ever binds — before any architectural change.

Beyond the headline comparison, the queued work is: multi-seed evaluation with confidence intervals (pilot probes are paired but single-seed); the flat-MLP-over-one-hots ablation if the CNN’s smeared single-

pixel features ever limit performance; a voltage-to-threshold actuation mapping to replace the current health-threshold abstraction; trailing (rather than frozen) position-delay sensing; grounding sweep densities to CMOS defect and calibration-drift rates instead of stress-test ranges; and, longest term, multi-droplet extension and physical validation of the sim-to-real gap — specifically, calibrating the abstracted health-scalar sensing model against measured capacitive readout noise, drift, and fault rates from fabricated hardware, without which the transfer risk of learned trust policies remains open. Each of these — scale-up, corner-cutting ablations, and hardware-grounded sensing calibration — is declared here as actively planned exploration, not background possibility.

References

- [1] Mahmoud Elfar, Yi-Chen Chang, Harrison Hao-Yu Ku, Tung-Che Liang, Krishnendu Chakrabarty, and Miroslav Pajic. Deep reinforcement learning-based approach for efficient and reliable droplet routing on MEDA biochips. *IEEE Transactions on Computer-Aided Design of Integrated Circuits and Systems*, 42(4):1212–1222, 2023. doi: 10.1109/TCAD.2022.3194808.
- [2] Gary Wang, Daniel Teng, and Shih-Kang Fan. Digital microfluidic operations on microelectrode dot array architecture. *IET Nanobiotechnology*, 5(4):152–160, 2011. doi: 10.1049/iet-nbt.2011.0018.
- [3] Basudev Saha, Bidyut Das, Vineeta Shukla, and Mukta Majumder. Pebble traversal-based fault detection and advanced reconfiguration technique for digital microfluidic biochips. *Journal of Electronic Testing*, 40(4):573–587, 2024. doi: 10.1007/s10836-024-06137-3.
- [4] Debraj Kundu, Gadikoyila Satya Vamsi, Karnati Vivek Veman, Gurram Mahidhar, and Sudip Kumar Roy. Reinforcement learning-based reliable synthesis of bioassays on MEDA digital microfluidic biochips. *ACM Transactions on Intelligent Systems and Technology*, 2025. doi: 10.1145/3773910.
- [5] Tung-Che Liang, Jin Zhou, Yun-Sheng Chan, Tsung-Yi Ho, Krishnendu Chakrabarty, and Chen-Yi Lee. Parallel droplet control in MEDA biochips using multi-agent reinforcement learning. In *Proceedings of the 38th International Conference on Machine Learning (ICML)*, volume 139 of *Proceedings of Machine Learning Research*, pages 6588–6599. PMLR, 2021.
- [6] Ling Zhang. Online diagnosis and self-recovery of faulty cells in daisy-chained MEDA biochips using functional actuation patterns. *Journal of Electronic Testing*, 2023. doi: 10.1007/s10836-023-06081-8.
- [7] Ling Zhang, Zipeng Li, Xing Huang, and Krishnendu Chakrabarty. Enhanced built-in self-diagnosis and self-repair techniques for daisy-chain design in MEDA digital microfluidic biochips. *IEEE Transactions on Computer-Aided Design of Integrated Circuits and Systems*, 42(10):3236–3249, 2023. doi: 10.1109/TCAD.2023.3244524.
- [8] Tanmoy Biswas, Pranab Roy, and Soumanetra Bose. Testing of MEDA-based biochip: A proposed technique for functional testing of symmetric set of modules. In *Proceedings of the 4th International Conference on Frontiers in Computing and Systems (COMSYS)*, pages 701–717, 2023.
- [9] Qinan Chen, Biao Liu, Chen Jiang, Handing Wang, Tsung-Yi Ho, and Bo Yuan. Reliable droplet routing in digital microfluidic biochips by reinforcement learning. In *Proceedings of the International Symposium on Electronics Design Automation (ISED)*, 2025.
- [10] Yuta Hamachiyo, Chiharu Shiro, Hiroki Nishikawa, Hiroyuki Tomiyama, and Shigeru Yamashita. Target-guided droplet routing on MEDA biochips considering shape-dependent velocity models and droplet splitting. *Micromachines*, 15(8):500, 2025.
- [11] Yuta Hamachiyo, Chiharu Shiro, Hiroki Nishikawa, Hiroyuki Tomiyama, and Shigeru Yamashita. Accelerated droplet routing on MEDA biochips considering shape-dependent velocity models through guidance and lower-bound constraints. *Applied Sciences*, 16(3):1421, 2026.
- [12] Tanmoy Biswas, Pranab Roy, and Krishnendu Basuli. MEDA based biochips: Technology, application and design automation issues – a brief review. *Integration*, 104:102436, 2025.

- [13] Ritwika Majumdar, Piyali Datta, Sagarika Chowdhury, and Rajat Kumar Pal. Advancement with digital microfluidic biochips towards sustainability and secured outcome: a comprehensive survey on specific design metrics. *Discover Electronics*, 1, 2024. doi: 10.1007/s44291-024-00018-x.
- [14] Pampa Howladar, Pranab Roy, and Hafizur Rahaman. MEDA based biochips: Detection, prevention and rectification techniques for cyberphysical attacks. *IEEE Transactions on Computational Biology and Bioinformatics*, 19(4):2345–2355, 2021. doi: 10.1109/TCBB.2021.3069380.
- [15] Tung-Che Liang, Zhanwei Zhong, Miroslav Pajic, and Krishnendu Chakrabarty. Extending the lifetime of MEDA biochips by selective sensing on microelectrodes. *IEEE Transactions on Computer-Aided Design of Integrated Circuits and Systems*, 39(11):3531–3543, 2020.
- [16] Hado van Hasselt, Arthur Guez, and David Silver. Deep reinforcement learning with double q-learning. In *Proceedings of the 30th AAAI Conference on Artificial Intelligence*, 2016. doi: 10.1609/aaai.v30i1.10295.
- [17] Andrew Y. Ng, Daishi Harada, and Stuart Russell. Policy invariance under reward transformations: Theory and application to reward shaping. In *Proceedings of the 16th International Conference on Machine Learning (ICML)*, pages 278–287, 1999.